

Name: Megha Samala

## Table of Contents

Project Direction Overview.....	2
Use Cases and Fields.....	2-7
Structural Database Rules.....	7-8
Conceptual Entity-Relationship Diagram.....	8
Full DBMS Physical ERD.....	9
Stored Procedure Execution and Explanations.....	9-11
Question Identification and Explanations.....	11
Query Executions and Explanations.....	12-13
Index Identification and Creations.....	13-14
Summary and Reflection.....	14-15

## Project Direction Overview

I would like to create a web application designed to track a user's media consumption across several different mediums. The application, which I will call "TunedIn," will be a place where users can log the different media they're currently consuming, have consumed, or plan to consume. The media in question can be movies, TV shows, books, music (in album form), and video games. Additionally, the user will be able to track information about the vendor of the media and use this information to organize and revisit their media buys as well as give recommendations to others

I believe an application like this is important because an overall media tracking system does not appear to exist, leaving people to log their consumption on multiple different websites dedicated to only a specific medium (think Letterboxd or MyAnimeList), or create a Twitter thread or blog dedicated to tracking, or alternatively not track their watches/reads/listens at all. As someone who considers herself to be scatterbrained, I also often jump from media to media without finishing a series or completing a playthrough, and I forget what episode I left on or what chapter of a book I was on, and sometimes I completely forget that I was in the middle of watching or reading something. When people ask me what I've been watching or reading or playing, or what my favorite movie or show is, I often blank and fail to answer the question. I can safely presume that I am not the only one with this issue, and therefore an application dedicated to tracking the status of someone's varied media consumption is a convenient solution. I also believe that tracking media consumption encourages users to think more critically about what they are consuming and how they feel about it, rather than finishing a series or book or game and simply moving on to the next one.

TunedIn will offer a place for users to input information about a piece of media, the main focus being on a name, a rating (if completed/desired), and a status on the consumption. Additional information can be added based on relevance, such as author information for a book, or the URL for something consumed online so it can easily be provided to a friend when recommending the media to them, or a link to another existing database-esque website which may contain a blurb/critic ratings of the media (again, think Letterboxd or MyAnimeList or IMDB, etc), and more. Users will be able to log what chapter of a book they are on, or what episode of a show they are on, so that if they drop it and choose to resume it later, they can easily pick up from where they left off. Additionally, users can track when they started and finished a piece of media, to log past watches, and/or determine if it might be time to revisit something they previously consumed and enjoyed. Each kind of media has information that is specific to it, as well as information that can be applicable to multiple kinds of media, such as a genre and a review. Through my multiple project iterations, I will be able to refine my scope and determine what information is most necessary and relevant for the user to have logged, and create the most useful version of my database.

## Use Cases and Fields

### Use Case #1: Account Setup

Because TunedIn will be a web application, it will be necessary to register to keep track of each individual user's media.

Field	What it stores	Why it's needed
username	This stores a username	Users can share usernames to

	associated with each account	see what another user is consuming, and users can also have multiple accounts if they would prefer
first_name	This stores the user's first name.	This can be displayed on screen and shared with other users.
last_name	This stores the user's first name.	This can be displayed on screen and shared with other users.
account_created	This stores the date on which the user created the account.	Users may want to know how long they have been using the application to track their media consumption.

### Use Case #2: Logging a Piece of Media

This is the supertype that users will log that contains common information across all media types before logging the specific subtype.

Field	What it stores	Why it's needed
media_name	This stores the name of the media.	This is necessary for the user to know what media they consumed.
date_released	This stores the date that the media was released.	This helps the user determine their viewing trends by date and also distinguish between potential media remakes or remasters with different release dates.

### Use Case #3: Logging a Movie

A movie will be one of the media types that the application will encourage users to log, and has fields that are specific to it. The other media types will track similar information and will be able to be differentiated in the database using a flag.

Field	What it stores	Why it's needed
director	This stores the name of the director of the movie.	This is useful for the user if they want to track how much they like movies by a certain director

		or want to sort by a certain director.
production_company	This stores the name of the production company of the movie, such as A24 or Studio Ghibli.	This is useful for the user to track how much they like the works of a certain company or sort by them.

#### Use Case #4: Logging a Genre

Since genres are applicable to any media type, users will be able to log them for any media consumed.

Field	What it stores	Why it's needed
category	This is a catch-all for the different media types where the category/large genres can be stored, such as things like fiction or nonfiction for books, documentary or short film for movies, RPG or shooter for video games, etc.	This allows for users to sort and search for certain pieces of media in a broad way.
main_genre	This stores the primary genre of a piece of media, which can differ in definition based on the user, but is generally a broader category such as horror, action, etc.	This allows for users to sort media by genre when organizing and also looking for previously logged media.
subgenre	This stores the subgenre of a piece of a media, which can be a specifier, like historical fantasy vs urban fantasy, etc.	This is optional but allows for more specific searching and sorting for the user.

#### Use Case #5: Logging a Review

Since reviews are applicable to any media type, users will be able to log them for any media consumed.

Field	What it stores	Why it's needed
out_of_ten	This stores a number rating out of 10.	This is a typical method of rating not just media, but anything, and will let users reflect on what they enjoyed or disliked.

review	This stores a text review.	This lets the user give clarification on the number rating and offers more information about the media.
review_link	This stores an optional link that leads to an external blog review or a website like Letterbox or Goodreads that the user may have posted their review on.	This lets the user organize relevant info in one place by allowing them to keep track of different websites they may have used to review a piece of media.

#### Use Case #6: Logging a Status

Each piece of media will have a consumption status that can be logged by the user.

Field	What it stores	Why it's needed
overall_status	This stores a general status like Planned, In Progress, Completed, Dropped, or On Hold.	This helps users keep track of all the media they want to watch, are watching, and have watched, as well as media they chose to not complete and media they plan to return to, which is handy for not forgetting their progress.
section	This stores a larger indicator of progress, such as a season of a TV show, a chapter for a book, a section/act of a video game, etc	This is helpful for giving context to the part of the media the user is on, or is an indicator that they finished it.
subsection	This stores a smaller indicator of progress, such as an episode of a TV show, a page of a book, a level of a video game, etc.	This is helpful for tracking the exact part of a piece of media that the user left off on so they can jump back in.
date_started	This stores the date when the user started consuming the media.	This is helpful for the user to know when they first began a piece of media and for tracking how long it is taking them to complete it.
date_finished	This stores the last date the user interacted with/finished the media.	This is helpful for tracking the amount of time the user has spent with the media and for

		tracking how long ago they interacted with the media to see if it is worth a revisit.
--	--	---

#### Use Case #7: Logging a Vendor

Media is offered by multiple vendors and these vendors can be frequented by users, and therefore are logged in the database.

Field	What it stores	Why it's needed
vendor_name	This stores the name of the vendor.	This is needed for the user to know what vendors they buyt from.
vendor_type	This stores the type of vendor, such as website, streaming service, or physical storefront.	This is useful for the user to know where they acquired a piece of media.
vendor_link	This stores a link to the vendor website, if available.	This can help the user track past visits to an online vendor and allow them to revisit it or recommend it.

#### Use Case #8: Vendor-Media Relationship

This connects the vendor to the piece of media and provides useful information about the purchase.

Field	What it stores	Why it's needed
vendor_name	This stores the name of the vendor and will reference the vendor table.	This lets users know where they bought or can buy a piece of media.
date_purchased	This stores the date of the purchase of the media from the vendor.	This can assist with tracking the consumption of media by time.
media_link	This stores a direct link to the media on the vendor site, if available/applicable	This can help the user revisit their purchase easier or recommend the purchase to a friend.
price	This stores the price of the media purchase, if applicable	This can help users budget and track their money spent on media and also play into their

		review of the media by determining if it was worth the price they paid for it.
--	--	--

#### Use Case #9: Series, Anthology, Collection

Some pieces of media are part of a series or anthology or collection, such as a book series or short film anthology, and users can specify which series a piece of media is part of.

Field	What it stores	Why it's needed
SAC_name	This stores the name of the series, anthology, collection or other extended body of work that a piece of media is part of.	This helps users see how their media is connected as well as track their consumption of a series or the like.
sub_SAC_name	This is optional and stores the name of a subseries/sub-extended body of work within a larger series.	This further helps users organize and search for media.
SAC_part	This is optional and stores a number to indicate which part of a chronological series the piece of media is (such as the second book in a trilogy, etc)	This helps users track completion of a connected/chronological body of media.

#### Use Case #10: Rating

Most pieces of media have a rating, and users will be able to log this.

Field	What it stores	Why it's needed
rating	This stores the rating of a piece of media, such as G or PG-13 or R in the case of movies, E or T in the case of video games, etc.	This allows users to sort by the rating of the media they are looking for and keep track of their consumption trends.

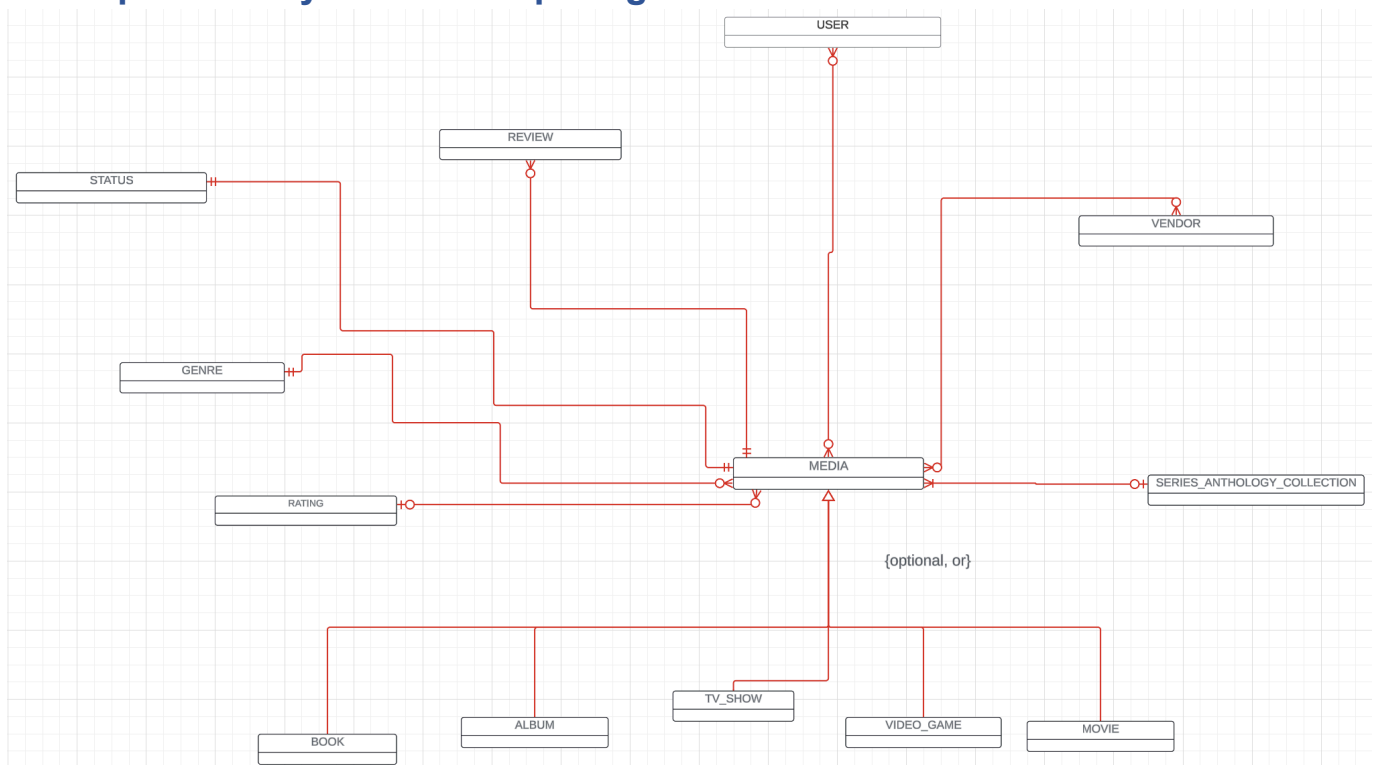
### Structural Database Rules

- 1) Each user may log zero to many pieces of media, and each piece of media may be logged by zero to many users.
- 2) A piece of media is a book, album, TV show, movie, video game, or none of these.
- 3) Each piece of media may have zero to many reviews, and each review will apply to one piece of media.

- 4) Each piece of media will have one status, and each status will apply to one to many pieces of media.
- 5) Each piece of media will have one genre, and each genre can apply to zero to many pieces of media.
- 6) Each piece of media can be part of one series/anthologies/collections, and each series/anthology/collection can contain many pieces of media.
- 7) Each piece of media can be purchased/rented from zero to many vendors, and each vendor can offer zero to many pieces of media.
- 8) Each piece of media may have one rating, and each rating can apply to zero to many pieces of media.

I updated rules 3 and 4 to reflect the feedback I received and the changed relationships between Status -> Media and Media -> Review.

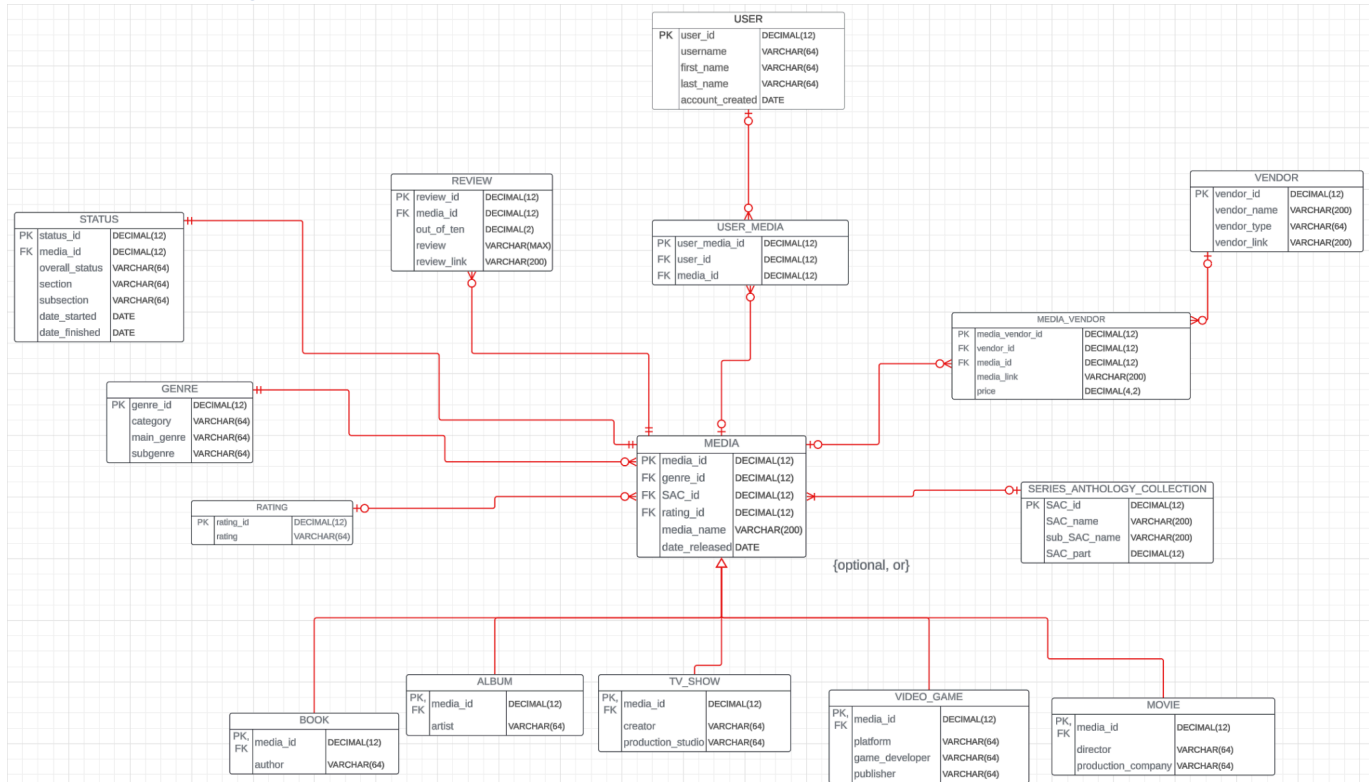
## Conceptual Entity-Relationship Diagram



My conceptual ERD has not changed much since my last iteration. Based on the feedback received, I changed the relationship between Media -> Review to one-to-many from one-to-one. I kept the relationship between Status -> Media as one-to-one because the information contained in each status entry is highly personalized. Because users can input the exact part they are on and when they started/finished a piece of media, it is unlikely that one status can be used by more than one user on more than one piece of media. Additionally, removing the status\_id foreign key from the Media entity makes it so that there don't have to be multiple entries of the same piece of media with the only difference being the status of each one.



## Full DBMS Physical ERD



I changed some of the relationships in my physical ERD as I did in the conceptual ERD, which is explained above. I also removed SAC\_id and genre\_id foreign key fields from my subtypes as these foreign keys are only supposed to be in the supertype, and I removed the date\_purchased field from media\_vendor to make it so that multiple media purchases aren't logged with the only difference between them being the date purchased, thus making the Media\_vendor entity more applicable to all users.

## Stored Procedure Execution and Explanations

```

143 -- adding users
144 CREATE OR REPLACE PROCEDURE add_user(
145     username_arg IN VARCHAR,
146     first_name_arg IN VARCHAR,
147     last_name_arg IN VARCHAR)
148     LANGUAGE plpgsql
149 AS
150 $$
151 BEGIN
152     INSERT INTO Users(user_id, username, first_name, last_name, account_created)
153     VALUES(nextval('user_seq'), username_arg, first_name_arg, last_name_arg, CURRENT_DATE);
154 END;
155 $$;
156
157 START TRANSACTION;
158 DO
159 $$ BEGIN
160     CALL add_user('msamala', 'Megha', 'Samala');
161     CALL add_user('nbardhan', 'Neha', 'Bardhan');
162     CALL add_user('aismail', 'Afra', 'Ismail');
163     CALL add_user('sanand', 'Soumitra', 'Anand');
164     CALL add_user('mkoleti', 'Manish', 'Koleti');
165     CALL add_user('tnguyen', 'Tim', 'Nguyen');
166 END $$;
167 COMMIT TRANSACTION;
  
```

```

169 -- adding rating
170 CREATE OR REPLACE PROCEDURE add_rating(
171     rating_arg IN VARCHAR)
172     LANGUAGE plpgsql
173 AS
174 $$
175 BEGIN
176     INSERT INTO Rating(rating_id, rating)
177     VALUES(nextval('rating_seq'), rating_arg);
178 END;
179 $$;
180
181 START TRANSACTION;
182 DO
183 $$ BEGIN
184     CALL add_rating('TV-Y');
185     CALL add_rating('TV-Y7');
186     CALL add_rating('TV-G');
187     CALL add_rating('TV-PG');
188     CALL add_rating('TV-14');
189     CALL add_rating('TV-MA');
190     CALL add_rating('G');
191     CALL add_rating('PG');
192     CALL add_rating('PG-13');
193     CALL add_rating('R');
194     CALL add_rating('NC-17');
195     CALL add_rating('E');
196     CALL add_rating('E10+');
197     CALL add_rating('T');
198     CALL add_rating('M');
199     CALL add_rating('NR');
200     CALL add_rating('Explicit');
201 END $$;
202 COMMIT TRANSACTION;

```

```

514 -- adding a media as a book
515 CREATE OR REPLACE PROCEDURE add_book(
516     media_name_arg IN VARCHAR,
517     date_released_arg IN DATE,
518     author_arg IN VARCHAR) LANGUAGE plpgsql
519 AS
520 $$
521 DECLARE
522     v_media_id DECIMAL(12);
523 BEGIN
524     SELECT media_id
525     INTO v_media_id
526     FROM Media
527     WHERE media_name = media_name_arg AND date_released = date_released_arg;
528
529     INSERT INTO Book(media_id, author)
530     VALUES(v_media_id, author_arg);
531 END;
532 $$;
533
534 START TRANSACTION;
535 DO
536 $$ BEGIN
537     CALL add_book('The Honjin Murders', CAST('04-AUG-2020' AS DATE), 'Seishi Yokomizo');
538     CALL add_book('Death on Gokumon Island', CAST('05-JUL-2022' AS DATE), 'Seishi Yokomizo');
539     CALL add_book('Mistborn: The Final Empire', CAST('17-JUL-2006' AS DATE), 'Brandon Sanderson');
540     CALL add_book('Jujutsu Kaisen', CAST('05-MAR-2018' AS DATE), 'Gege Akutami');
541     CALL add_book('Bleach', CAST('07-AUG-2001' AS DATE), 'Tite Kubo');
542 END $$;
543 COMMIT TRANSACTION;

```

Pictured above are some of the stored procedures I have implemented in my project to assist with populating the database tables. I have a variety of stored procedures to represent different use cases that I have outlined for my project, such as a procedure for adding a new user account, a procedure for adding content ratings to select from, and more, and the comments above each procedure talk about the purpose of the procedure. My database is populated using the stored procedures as well as manual inserts. All of the stored procedures can be viewed in my SQL script file. Additionally, I added error checking functionality, as shown below:

```
169 -- throw error if username is already taken when creating new user account
170 CREATE OR REPLACE FUNCTION user_username_func()
171 RETURNS TRIGGER LANGUAGE plpgsql
172 AS $$
173 BEGIN
174     IF EXISTS(
175         SELECT username FROM Users WHERE username = NEW.username
176     ) THEN
177         RAISE EXCEPTION USING MESSAGE = 'Username is already taken',
178         ERRCODE = 23000;
179     END IF;
180     RETURN NEW;
181 END;
182 $$;
183 CREATE TRIGGER user_username_trg
184 BEFORE UPDATE OR INSERT ON Users
185 FOR EACH ROW
186 EXECUTE PROCEDURE user_username_func();
187
188 INSERT INTO Users(user_id, username, first_name, last_name, account_created)
189 VALUES(nextval('user_seq'), 'msamala', 'Sdnfsdf', 'Hjsdfjb', CURRENT_DATE);
190
```

Data Output	Messages	Notifications
ERROR: Username is already taken CONTEXT: PL/pgSQL function user_username_func() line 6 at RAISE		
SQL state: 23000		

This function makes it such that usernames cannot be reused for different accounts.

## Question Identification and Explanations

First Query: Which subscription service is the most used for watching TV shows? This question is useful because it allows users to determine which subscription service appears to have the most media available, and is therefore the most worth paying for compared to the others. Additionally, if a subscription service is being used less, users can be made aware and can choose to cancel their subscription.

Second Query: What is the best reviewed video game that falls under the “fantasy” genre? This question is useful because it is a good example of the kinds of questions users will ask while using the database to determine what piece of media they might choose to interact with next. Creating a query for this question allows for users to see what high quality media, specifically video games, is out there being recommended to them in a genre of their choice.

Third Query: What is the average rating given to a piece of media by each user? This question is useful because users can see which users are harsher critics than others, and use the average rating given by each user to determine which users are consuming more quality media, and then use that information to determine what media they might interact with next - they may take recommendations from users that give a higher rating to media on average.

## Query Executions and Explanations

First Query:

```
908 -- first query: Which subscription service is the most used for watching TV shows?
909 SELECT vendor_name, COUNT(*) AS use_count
910 FROM Vendor
911 JOIN Media_vendor ON Media_vendor.vendor_id = Vendor.vendor_id
912 JOIN Media ON Media.media_id = Media_vendor.media_id
913 JOIN TV_show ON TV_show.media_id = Media.media_id
914 GROUP BY vendor_name;
915
```

Data Output		Messages	Notifications
	vendor_name character varying (200)	use_count bigint	
1	Netflix	3	
2	HBO Max	1	

The query that answers the first question has 4 tables joined: Vendor, Media\_vendor, Media, and TV\_show, and uses the COUNT function to count how many TV shows were watched on each streaming subscription service. From this query, it is clear that Netflix is the most used service for watching TV shows.

Second Query:

```
941 -- second query: What is the best reviewed video game that falls under the "fantasy" genre?
942 SELECT media_name, out_of_ten
943 FROM Review
944 JOIN Media ON Media.media_id = Review.media_id
945 JOIN Genre ON Genre.genre_id = Media.genre_id
946 JOIN Video_game on Video_game.media_id = Media.media_id
947 WHERE main_genre = 'Fantasy' OR subgenre = 'Fantasy'
948 ORDER BY out_of_ten DESC;
949
```

Data Output		Messages	Notifications
	media_name character varying (200)	out_of_ten numeric (2)	
1	Final Fantasy VII Remake	9	
2	Final Fantasy XVI	9	
3	Fire Emblem: Three Houses	8	

This query joins the Video\_game subtype with the Media supertype and the Review and Genre tables to answer the question. The joins combined with the WHERE clause limits the results to video games that only have “Fantasy” listed as one of the genres, and the ORDER BY clause puts the highest rated video game at the top of the result set. From this query, it is shown that Final Fantasy VII Remake and Final Fantasy XVI are tied for the highest rated fantasy video games in the database.

Third Query:

```

950 -- third query: What is the average rating given to a piece of media by each user?
951 CREATE OR REPLACE VIEW User_ratings AS
952 SELECT username, media_name, out_of_ten
953     FROM Review
954     JOIN Media ON Media.media_id = Review.media_id
955     JOIN User_media ON User_media.media_id = Media.media_id
956     RIGHT JOIN Users ON Users.user_id = User_media.user_id
957     ORDER BY username, out_of_ten DESC;
958
959 SELECT user_ratings.username, ROUND(AVG(out_of_ten), 2) AS avg_rating
960 FROM User_ratings
961 GROUP BY user_ratings.username;
962

```

Data Output
Messages
Notifications

	username character varying (64)	avg_rating numeric
1	aismail	8.25
2	mkoleti	[null]
3	msamala	8.40
4	nbardhan	9.00
5	sanand	8.00
6	tnguyen	7.00

The view I created is called User\_ratings and shows all of the ratings out of 10 each user has given to pieces of media, and includes all users regardless of if they have reviewed any media or not by using the RIGHT JOIN on the Users table. The ORDER BY clause displays the results in alphabetical order of usernames, and descending order of ratings out of 10 within each username. Using the view in my query, I used the AVG aggregate function to retrieve the average rating given by each username to a piece of media. From this we see that some users give lower ratings on average than others, and one user has not rated any media yet.

## Index Identification and Creations

The primary keys of my database are already indexed, and the foreign keys of my database need to be indexed. I have concluded that most of the foreign keys do not need unique indexes because most of them are non-unique - for example, foreign keys such as genre\_id and SAC\_id, among others, do not need unique indexes because many pieces of media can fall under the same genre or series/anthology/collection. The only foreign key that needs to have a unique index is Status.media\_id because of the 1:1 relationship between Status and Media. Additionally, in one of my queries I used main\_genre and subgenre in the WHERE clause, and therefore will be giving non-unique indexes to these two columns in the Genre table. Below is the creation of the indexes:

```

946 -- index creation
947 CREATE INDEX media_genre_idx
948 ON Media(genre_id);
949 CREATE INDEX media_SAC_idx
950 ON Media(SAC_id);
951 CREATE INDEX media_rating_idx
952 ON Media(rating_id);
953 CREATE INDEX user_media_user_idx
954 ON User_media(user_id);
955 CREATE INDEX user_media_idx
956 ON User_media(media_id);
957 CREATE INDEX media_vendor_media_idx
958 ON Media_vendor(media_id);
959 CREATE INDEX media_vendor_vendor_idx
960 ON Media_vendor(vendor_id);
961 CREATE INDEX review_media_idx
962 ON Review(media_id);
963 CREATE UNIQUE INDEX status_media_idx
964 ON Status(media_id);
965 CREATE INDEX book_media_idx
966 ON Book(media_id);
967 CREATE INDEX album_media_idx
968 ON Album(media_id);
969 CREATE INDEX tv_show_media_idx
970 ON TV_show(media_id);
971 CREATE INDEX video_game_media_idx
972 ON Video_game(media_id);
973 CREATE INDEX movie_media_idx
974 ON Movie(media_id);
975

```

Data Output Messages Notifications

CREATE INDEX

Query returned successfully in 175 msec.

```

976 -- query-driven indexes
977 CREATE INDEX main_genre_idx
978 ON Genre(main_genre);
979 CREATE INDEX main_genre_idx
980 ON Genre(subgenre);
981

```

Data Output Messages Notifications

CREATE INDEX

Query returned successfully in 175 msec.

## Summary and Reflection

My database is for a web application called “TunedIn” in which users can track their media consumption across various mediums, as well as media purchases from various vendors. The application allows users to rate their media and track their progress in consuming it.

This week I made a few tweaks to my physical and conceptual ERDS as well as my structural database rules based on feedback from last week. The main progress this week was creating a large amount of stored procedures for my database and populating the database using these procedures and a large number of manual inserts. Additionally, I created queries based on relevant questions to the database, and indexes to make querying more efficient. All of these, including parts I could not include in this iteration document can be viewed in my SQL script file.

Ultimately, I am excited to wrap up the project, but I had a few doubts regarding my overall database structure while I was building it out, mainly realizing that things like Status and Reviews are personal to each user but I cannot think of how to link these to the User at this stage. I appreciate any feedback regarding this matter and any other part of this week's iteration.